

LeadsTo Software

Lourens van der Meij
E-mail: `lourenstcc@gmail.com`

June 20, 2016

Abstract

This document describes the LeadsTo software in detail. It starts out as an investigation into details of the algorithm.

Contents

1	LeadsTo core design and implementation	3
1.1	Introduction	3
1.2	run_simulation/2	3
1.3	LeadsTo specifications	3
1.3.1	Sources for information	3
2	Details	3
2.1	Times: Handled Time, Setup Time, Start Time, End Time	3
2.2	Traces	4
2.2.1	Datastructures	4
2.2.2	Algorithm Variables	5
2.2.3	Loading and Trace Generation	5
2.2.4	Saving traces	7
3	Details	8
3.1	Constants	8
3.2	Model Checking	8
3.3	recwait	8
3.4	Following the progress of leadsto rules	8
4	Working backwards	9
4.1	Leadsto times e, f, g, h	10
4.2	Invariant	10
4.3	Sketch of the algorithm	10
4.4	Rule States	10
4.5	Garbage collection	10
4.6	Runtime Algorithm Predicates	11
5	wait entries	13
5.1	wait_var	14
6	So far	14
7	"Syntax" of Leadsto Specifications	14
8	TODO	16
8.1	waitvar, waitfired	16
8.1.1	Debugging, verifying waitvar FV,FVL aspects	16
8.2	Solving bugje-nondeterminism	17
8.3	analyse case	17

1 LeadsTo core design and implementation

1.1 Introduction

A leadsto specification is encoded as a pl file containing predicates.

1.2 run_simulation/2

Here we describe the main leadsto procedure `run_simulation/2`, `run_simulation(File, Frame)`. This leadsto specification *File* is loaded into the leadsto runtime:

1. The predicates are loaded into module `spec`.
2. After that, all terms in the input file are preprocessed, often leading to asserted dynamic predicates in the current(algo) module.

There seems to be almost no compilation at this stage and it looks like terms in the `spec` module often are asserted as facts into `algo` without any transformation.

Some translation of sortdefs is performed. If a sort contains less than 100 ground terms, it is instantiated, otherwise the sort definition is left as is. This leads to `spec:sortdef(Sort, Terms)`. The source contains a beginning of a new way of encoding sort definitions.

3. If the specification contained a model specification, we run each model instance after setting up model parameters. Otherwise we perform a single run.
4. Running the specification
 - (a) setup of the runtime (first part of `runspec1/0`)
 - (b) performing the firing of rules (`runspec_rest/0`)
5. Saving the generated trace

1.3 LeadsTo specifications

1.3.1 Sources for information

The file `userman.html` contains the documentation for most allowed constructs in LeadsTo specifications.

The file `olddoc/syntax.txt` also describes the syntax. See section 7.

2 Details

2.1 Times: Handled Time, Setup Time, Start Time, End Time

setuptime, `TSetup` In practice identical to `TStartup`, but there are options for defining `TSetup` to have a value before `TStart` so that leadsto rules could fire for antecedent values before

TStart and use `cwa (Atom)` derived values to make them fire. The current value of TSetup is stored in `dyn_setup_time (TSetup)`.¹

handledtime Handled time is initialized by `setup_unknown_or_cwa/2` to TStart. There is an additional implied condition on handled time: "You should never ask for values before TSetup" and "All atoms that have no explicit trace entry before THandled have value false if cwa, unknown otherwise".

starttime, TStart Start time. The algorithm uses TSetup, but TStart still plays a role, when storing traces, only values at/after TStart are saved.² TStart is stored as `dyn_start_time (TStart)`, but only called through `start_time/1`.

endtime If not specified there currently is a default of 200 (see `end_time/1`)

TSetup and TSetup are set up in `do_setup_time (TStart, TSetup)`. They may contain specification constants.³

HandledTime is incremented in `handled_time_step/1` and `runspec_rest/0` ensures that at the end of the leadsto algorithm `HandledTime >= Endtime`.

2.2 Traces

2.2.1 Datastructures

Traces are stored as Prolog facts, each fact represents values of a single ground atom.

Values of an atom over time are represented as a list:

```
[range(23.0, 24.0, true), range(17, 18, true)]
```

The ranges are ordered, the latest time range first. During the execution of our leadsto algorithm, only necessary values are stored, unknown ranges or false ranges where `cwa (Atom)` holds are not part of the trace.

Traces are generated by the main algorithm in module `algo`. They are internally stored as `dyn_atom_trace (AtomKey, Atom, AtomTrace)` facts. For performance reasoning traces that can no longer play a role in the algorithm are backed up into `dyn_atom_trace_backup/3` facts.⁴

In saved traces all values are represented. Saved traces will only contain atom values in the range `start_time` to `endtime`. The leadsto algorithm may derive values outside of that range.⁵

¹ We had command line options `setup_maxg` and `setup_maxfg` for that purpose that would introduce TSetup based on the maximum leadsto rule values for *f* and *g*. But the sourcecode says this is not supported. So, in practice `TStart == TSetup`

²TODO:check this!

³In `do_setup_times/2` the values are passed through `tr_basic_element (Term, [], TermOut)` that will substitute `spec:constant (Name, Val)` occurrences. See section 3.1.

⁴TODO:I seem to remember that at places in the algorithm we depend on there either being `dyn_atom_trace/3` or `dyn_atom_trace_backup/3`.

⁵ TODO: We should discuss alternatives:

1. Why not save only necessary values in saved traces?
2. If saving everything, why not compact the timerange? `[range(T1, T2, TFU1), range(T2, T3, TFU3), ...]` into `[T1-TFU1, T2-TFU2, .. TE-[]]`

`filled_atom_trace/3` and `atom_trace/3` are probably only used for saving and printing out results.

`dyn_atom_trace`(*AtomKey*, *Atoma*, *AtomTrace*)

`dyn_atom_trace_backup`(*AtomKey*, *Atoma*, *AtomTrace*)

`atom_trace`(*AtomKey*, *Atoma*, *AtomTrace*)

Looks for both `dyn_atom_trace/3` facts and `dyn_atom_trace_backup/3` facts.

`filled_atom_trace`(*AtomKey*, *Atoma*, *AtomTrace*)

Same as `atom_trace/3`, but fills in cwa values up to `EndTime`

`find_atom_trace`(*?Atom*, *+AtomTrace*)

Only looks for `dyn_atom_trace/3` facts, not `dyn_atom_trace_backup/3`.

`atom_key`(*+Atom*, *?AtomKey*)

Returns *AtomKey* given *Atom*. *Atom* must be ground. The predicate defines the relationship by means of `term_to_atom(Atom, AtomKey)`; there are problems with this way of defining the hash-value. I seem to remember that floating point numbers could give wrong results?

The *AtomKey* - *Atoma* pairs

2.2.2 Algorithm Variables

`dyn_sim_status`(*File*, *Status*)

says in what stage of loading and running the algo algorithm is. *Status* is loaded, running or done.

`dyn_currently_loaded`(*Kind*, *File*)

says what *File* is loaded and what *Kind*, where *Kind* is `trace` or `sim`.

2.2.3 Loading and Trace Generation

`load_simulation`(*+File*)

The specification *File* is loaded into module `spec`. The source code seems somewhat complex: The module `spec` is set up, `discontinuous/1` directives are generated for all leadsto specification terms, the terms are read from *File* and asserted one by one into the `spec` module.

Command line constants are added to module `spec` (see section 3.1).

Then, the leadsto specification is read one more time, and each `Term` is passed on to `handle_term/1`. Most terms are handled by asserting dynamic facts into module `algo`. Some of those are 1-1 translations, others are not.

`model(Model)` is translated into `dyn_model(Model)`, after checking that there is only one such term.

`cwa(F/A)` is translated into `dyn_cwa(FunctorTerm)`.

The most interesting things happen with interval leadsto specification terms and leadsto rule terms. Interval rules are converted into two standard forms and stored by `initialise_interval/3` and `initialise_interval_periodic/4`.

`handle_term/1` itself simply asserts leadsto rules into `dyn_leadsto(RuleId, Vars, Antecedent, Consequent, Delay)` facts. `setup_leadsto/6` processes these facts further in `runspec1/0` at algorithm startup.

Finally `update_sorts/0` performs some pre compilation of sort definitions.⁶

initialise_interval(+Range, +Vars, +LiteralConjunction)

Handles initial setup of all non-periodic interval rules. In this phase the predicate asserts `dyn_interval(i(Range, Vars, LiteralConjunction))`. `setup_rt_intervals/0` processes the terms further at algorithm startup.

initialise_interval_periodic(+Range, +Period, +Vars, +LiteralConjunction)

the predicate asserts `dyn_interval(i(Range, Period, Vars, LiteralConjunction))`. `setup_rt_intervals/0` processes the terms further in `runspec1/0` at algorithm startup. `initialise_interval/3` and `initialise_interval_periodic/4` deal with them.

Leadsto rule terms are translated into `dyn_leadsto(I, Vars, LitDisConj, AndLiterals, Delay)` facts.

`load_simulation/1` sets `dyn_sim_status(File, loaded)`.

reset_sim_info

clears the content of spec together with other run time information.

runshowspec(+Frame)

Two parts, `runspecdo/1` and `show_results/2`.

runspecdo(+Frame)

(Functionality in `runspec/1`). If we are dealing with a model, we initialize the output trace common to all model traces, then for each model instantiation we call `runmodel/4` that does `runspec1/0` and cleans up after itself for the next `runmodel/4`.

If there is no `model/1`, we call `runspec1/0` followed by `savetrace/1`.

runspec1

This procedure calls `do_setup_time/2` that sets up `TSetup` and `TStart`.

In `runspec/1` we perform: `setup_rt_intervals/0`, `setup_unknown_or_cwa/2`, `setup_leadsto/6` for each leadsto rule `dyn_leadsto/5`,

⁶ TODO: It is unclear what happens to sortdefs at this stage. Questions are:

- What terms are used? Are they `spec:sortdef/2` and `spec:sortdef/4`?
- What is the role of `dyn_sortdef/5`?
- Are constants somehow substituted into sortdef elements?

Question surrounding `load_simulation/1`:

- Why is the leadsto specification scanned twice. The terms will be probably be walked through even one more time to set up the algorithm.

get_model_checking_p_rules/0 (??), setup_atom_state_boundaries/0
and finally do the real reasoning in runspec_rest/0.

setup_rt_intervals

For every `dyn_interval/1` term we perform `init_interval_callbacks/10` where all but the first three arguments are callback variables or callback predicates.

The setting of the interval rules does some detailed steps such as variable instantiation. Finally this leads to changes to `dyn_atom_trace/3`. See

```
initialise_interval_p(Interval, P, Vars, Form1) :-
    (
        instantiate_vars(Vars, VarsInst),
        tr_range(Interval, VarsInst, T1, T2),
        tr_basic_element(P, VarsInst, P1),
        ...
    )
```

setup_leadsto(+TStart, +Vars, +Antecedent, +Consequent, +Delay, +RId)

(Called by `runspec1/0`). It translates and transforms the rule using `init_interval_callbacks/10`, then calls `setup_lt_internalL/5` which simply calls `setup_lt_internal/6`. That predicate simplifies Antecedent and Consequent terms using `simplify_term/4`. In some cases the rules simplify to specifying constant Antecedents. Other cases are passed on to `setup_nontrivial_leadsto/5`. It stores the compiled leadsto rule as `dyn_lt_rule(Id, AnteLits, ConseRId, PVOutC, Delay, RId)` but also does initial firing and instantiation the leadsto rule data structures by calling `setup_lt/6` which calls `setup_lt_normed/8` with argument `Removed = initial`, which is part of the run time algorithm.

init_interval_callbacks(TmInf, Vars, Forms, TmInf1, Vars1, Forms2, InvldVars, InvldTimeInfo, ActPreInstttiated)

is used for setting up interval rules and leadsto rules. It instantiates variables, also takes care of `forall/2` terms(instantiates them).

2.2.4 Saving traces

Traces are saved in two stages by

savetracesetup(+File, +Frame, -Telling)

Saves constants and sets up trace storage stream.

savetrace1(+TraceName)

Saves the trace itself. (If *TraceName* is `[]`, trace will not have trace id.)

savemodelspec_cleanup(-TellStream, +ModelInstanceTraceName)

If the *leadsto specification* contains a model, the separate model instances saved.

3 Details

3.1 Constants

One can define *specification constants* constants that will be substituted into leadsto specification elements. Within a leadsto specification we use:

```
constant(Name, Value).
```

From the command line one can specify `-constant Name=Value`. This adds a constant to the specification. Value must be a valid ground Prolog term.

`set_option_constant/1` handles this by asserting `dyn_add_cmd_constant/2`.

`util:load_cmd_constants/0` loads those constants into module `spec` as `constant(Name, Value)` facts.

Constants are substituted by the procedure `tr_basic_element(Term, [], TermOut)`. Constants are stored as `spec:constant(Name, Val)` values.

3.2 Model Checking

The source contains code labelled *model checking*. I do not remember whether this code ever worked. I seem to remember I tried converting the leadsto model into some state based form.

Makefile contains an example call of using modelchecking:

```
./leadsto -local -modelchecking spec/heartn.lt
```

The only visible result seems to be some debugging info on the screen.

A first look at the code in *modelchecking.pl* does not make anything clear yet.

There is a document `olddocs/modelchecking.doc` that may provide background. I fear that the code that is still present in `algo.pl` never really did anything.

3.3 recwait

Within `algo` the two choices for representing algorithm state are mixed too much with the rest of the code. `recwait/0` is the switch between storage as recorded and storage as a dynamic clause. Sometimes code seems to be copy/pasted. But, it seems that backtracking in the recorded database and backtracking in the asserted database works differently, see `update_activity_times1/1`.

3.4 Following the progress of leadsto rules

We start with `setup_leadsto/6` where the arguments are almost identical to the values in the Leadsto specification.

setup_leadsto(TStart, Vars, LitDisConj, AndLiterals, Delay, RId)

where the arguments are almost identical to the values in the Leadsto specification. Then `init_interval_callbacks/9` transforms some constructs such as forall.

After a number of steps involving normalizing conjunctions and disjunctions⁷ and partial evaluation pruning out true and false results, `setup_lt/6` is called.

⁷TODO:Verify whether disjunction is allowed

setup_lt(Ante, Conse, Vars, Delay, Id, RId)

The encoding of the antecedent is responsible for generating code. If a Term is a comparison operator, code is generated for that, if a term is an arithmetic expression, code is also generated.

We pass on some (incomplete) data structures within `setup_lt/6`. `code_form/4` uses `ds_d(AnteResult, VarsIn, PVIn)` and `ds_d(AnteTail, VarsOut, PVOut)`. `AnteResult` is a difference list. Therefore often `AnteTail` is set to `[]`.

`setup_lt/6` calls `setup_lt_normed/8`. The result is stored as `dyn_lt_rule(Id, AnteLits, ConseRId, PVOutC, Delay, RId)` but more important, `setup_lt_normed/8` is called, it leads to `setup_lt_wait_var/12`.

code_form(+AnteConse, +PosNeg:[pos, neg], DIn, DOut)

`code_form/4` is used for `Ante` and `Conse`.

Each `AnteConseTerm` is translated as a list element in `AnteResult`. `L = ds_litd(Atom, PosNeg, PreOps, PostOps, PostConds)` where `Atom` can be true or any other value. It seems that its translated value is not tested in any way.

Within `code_form/4`, `tr_arg_prolog1(Term, PVIn, Term1, Inst, DSTAIn, DSTAOut)` is used where

```
DSTAIn = ds_ta(VIn, PVIn, [], [], []),  
...  
DSTAOut = ds_ta(VOut, PVOut, PreOpsOut, PostOpsOut, PostCondsOut),
```

`Inst` should result in `Inst == inst`.

tr_arg_prolog1(Term, PVIn, Term1, Inst, DSTAIn, DSTAOut)

`tr_arg_prolog1/6` translates `leadsto` variables into Prolog variables, their relationship is stored and retrieved in `PVIn/PVOut`, by `var_pl_to_var_list/6` and `var_pl_from_var_list/5`.

The first encounter of a `leadsto` variable in a `code_form/4` has `Inst = next`, a later one gets `Inst = inst`. `Inst` values can be `inst`, `next`, `var`, `mixed`.

`tr_arg_prolog1/6` is also responsible for substituting *spec_constants*.⁸

code_conse(Conse, VOut, PVOut, Id, ConseRId, PVOutC)

translates `Conse` through `code_form/4`, but true `ConseLits` are removed.⁹

The consequent is encoded as `ds_cr(ConseLits, ds_ri(Id))`, but `pxor` consequents are treated differently.¹⁰

TODO: Looking at the code, it seems that we do not reorder the `AnteLiterals` depending on intermediate results.

4 Working backwards

Meaning, trying to reconstruct the algorithm from the start.

⁸TODO: Check whether Atoms could end up as Prolog variables, look at `code_atom/4` where `Inst` is ignored in the code. Can `Inst` be `var` or `mixed` there?

⁹(TODO: why not from `Ante`?).

¹⁰TODO

4.1 Leadsto times e, f, g, h

We limited $e, f, g, h: e, f, g, h \geq 0$ and if $h == 0$ then g must be 0. But also, $e + f + g + h > 0$.¹¹

$e, f, g, h:$

Once an antecedent holds for duration $g + T$, a delay is set between e and f , and the antecedent will hold during $h + T$. So, even if a rule has fired, we need to remember that it has fired and as long as the antecedent may continue to hold, the consequent will be propagated for a longer time.

4.2 Invariant

`HandledTime`: Everything that can be derived, has been derived for $T \leq \text{HandledTime}$. CWA atom values do not have to be instantiated, probably will not be instantiated to false values.

4.3 Sketch of the algorithm

All rules that could still fire are inspected, their antecedent effect is exhaustively tested up to `HandledTime` at least.

After everything has fired, we inspect all waiting antecedents, and look at time their first result could come in. And the minimum value becomes the next `HandledTime`, unless this minimum value is not after `HandledTime` (could it be smaller?). It looks like the algorithm currently simply gives up if there is a rule that could fire at `HandledTime`.

It is probably important that together with setting `HandledTime`, every rule that has some continuation has its effect propagated till the new `HandledTime`.

This would make the invariant more precise: Every rule has its state updated in such a way that the antecedents have been checked up to the new `HandledTime`.

4.4 Rule States

Rules can contain variables, that is, antecedent literals can have variables. There can be more than one separate state per rule.

It could be that the first N literals of the antecedent with some specific instantiation are valid in some time range $T1 - T2$.

We will look strictly left to right.

But, the extending of the fired rules is done in reverse, why? Probably because we wish to extend the range as far as possible.

4.5 Garbage collection

Is complex and not documented. In the source some explanation is given, `olddocs/bugdev.txt` also contains some explanation.

`dyn_handled_wait_var_instance/4` is complex, asserted in `mark_handled/7`. `mark_handled/7` is only called in `rm_gc_wait_vars1/8`.

In `fail_filter_handle/15` a comment says "ignoring handled instance" and then indeed skips handling the FV instantiation.

In `instantiate_op/16` we skip `setup_lt_normed/8` if we encountered the matching `dyn_handled_wait_var_instance/4`.

¹¹Why those requirements? We probably do not want to reason without delay.

Another fact: `dyn_rm_gc_wait_vars/8`. But that seems clear: assertion takes place only in `rm_gc_wait_vars/8` and depends on the `postpone_rm_gc` option being set. The entries are handled by `postponed_rm_gc/0` where indeed the postponed entries are handled by `rm_gc_wait_vars1/8`.

`postponed_rm_gc/0` is called at the end of `update_activity_times/1`. This seems to be harmless?

We should at least make sure that a duplicate handling of a `wait_var` entry does not lead to erroneous results. That could be the case if we would derive another delay if $e < f$. We need to be able to detect such situations!

Lets focus on `rm_gc_wait_vars/8`. Called by `rm_gc_wait_vars_th/10`, which is called only by `get_new_tholds/15` when the currently inspected *AnteHolds* `ds_lh` term fails from *Tholds*. Also called by `setup_lt_wait_g/12`:

- called by `update_lits_fired/8` but preceded by an implementation error.
- by `update_new_true_range/7` when from `Tlo` the current Literal is blank. If there the current Literal has fail value, `update_retrace1/12` is called.
- called by `update_retrace1/12`.

Probably `rm_gc_wait_vars/8` ensures ok result after a wait entry was removed?

`rm_gc_wait_vars`(*Atom, PN, Id, IdTerm, AnteHolds, AnteHoldsDone, ToDoAnte, Removed*)

updates `wait_var` situation. We wish to remove the FV as handled entry. It delegates the functionality to `rm_gc_wait_vars/8` or postpones garbage collection.

First `rm_wait_var/14` is called that simply retracts the entry. Then the FV value we are dealing with now is removed from FVL and then the `set_wait_var/14` is called with the updated content. TODO: ok, here we have a problem, the removal of the instantiation is linked to the new time, whereas the existing wait entry holds for an earlier time...?

Then `mark_handled/7` is called.

Then, after instantiating the wait entry with the current FV, we remove more, using `rm_gc_lit_data_deps_new/7`.

4.6 Runtime Algorithm Predicates

`setup_lt_normed`(*AnteTODO, AnteHolds, TMin, THolds, ConseRId, PV, Delay, Removed*)

We know *AnteHolds* holds in range *TMin, THolds*, we need to continue with *AnteTODO*. *AnteHolds* has all Prolog variables instantiated and excluded from FVL `wait_var` occurrences.

If *AnteTODO* becomes [],

`setup_lt_conse`(*AnteHolds, TMin, THolds, ConseRId, Delay, Removed*) is called. Otherwise, `setup_lt_notground/9` is called which calls `setup_lt_notground.fv/13` after setting up the free variable arguments.

`setup_lt_conse`(*AnteHolds, TMin, THolds, ConseRId, Delay, Removed*)

We know The whole antecedent holds between *TMin* and *THolds*. If $THolds \geq TMin + G$, we calculate $T3$ is $TMin + G + Delay$ and $T4$ is $THolds + Delay + H$ and then call `schedule_fire`(*ConseRId, T3, T4*) which sets

`dyn_schedule_fire(ConseRId, T3, T4).`¹² They are fired by the repeat, set_state, handle_fired sequence in `runspec_rest/0`.¹³ After `schedule_fire/3` we do `setup_lt_wait_fired/6`. If `THolds < TMin + G` we call `setup_lt_wait_true/5`.

`setup_lt_wait_fired/6` stores a `wait_fired/5` fact.¹⁴

setup_lt_notground_fv(+TStart, +FV, +FVL, +LitData, +ToDoAnte, +AnteHolds, +THolds, +ConseRId, +PV, +Delay, -AnteHolds holds for Time Interval between *TStart* and *THolds*. *ToDoAnte* is the conjunction that needs to hold. *LitData* is the Literal under investigation. *FV* are the free variables of the Literal and *FVL* is a list of instantiations that have been dealt with elsewhere.

Delay is `efgh(E, F, G, H)`, *ConseRId* is the consequent.

Removed indicates the source of the call. In case of `update_activity_time1(wait_var...)`, the `wait_var` term is Removed and the Removed is propagated along.

PV is probably the characterization of the variables: `pv(Arga, Sorta, Kinda, Argla)`.

FV, *fv* are abbreviations of Free Variables (Prolog variables).

`setup_lt_notground_fv/13` is called by `setup_lt_notground/9`, it by `setup_lt_normed/8`, it by `setup_lt/6`. `setup_lt/6` is called from `setup_nontrivial_leadsto/5`. `setup_lt_normed/8` is also called by `instantiate_op/16` (called in `add_default_cwa/17`) and `fail_filter.handleRR/16`, part of `filter_defaults.handle_others/14` in `setup_lt_notground_fv/13`.

We probably handle the rule in this call up to `HandledTime`.

Now, if at the call `THolds < HandledTime`, we start all over for this partially instantiated sequence of literals by calling `get_new_tholds/15`. Apparently the order of literals is reversed here. `get_new_tholds/15` is called here and by `get_new_tholds/15` itself.

If `THolds >= HandledTime` `setup_lt_notground_fv/13` analyses the current *LitData* first for all *AtomTraces* in `filter_defaults.handle_others/4` and all *cwa* matches in `add_default_cwa/17`, handles those separately by `check_fire_isolated/5` or `setup_lt_normed/8`. All those instantiations of *FV* handled separately are added to *FVL* and we finally call `setup_lt_notground1default/12` that calls `setup_lt_wait_var/12` with *Atom*, *FV*, *FVL* values implying that *FV* from *FVL* has been handled, but we need to check other instantiations.

get_new_tholds(AnteHoldsTODO, AHDone, TStart, THoldsNew1, Tholds, FV, FVL, LitData, ToDoAnte, ConseRId, PV, Del
Called by `setup_lt_notground_fv/13` and by `get_new_tholds/15` itself.

AnteHoldsTODO is an earlier instantiated sequence of literals that needs to be extended in range up to *THandled* (or further?). *AHDone* is the sequence that has been checked and hold between *TStart* and *THoldsNew1*. *Tholds* is the result.¹⁵

¹²I left out details dealing with `pxor` aspects. What is the reason for `schedule_fire/3`, why postpone?

¹³Why not fire immediately?

¹⁴It is confusing that two implementations of this waiting are present in the code, depending on the `recwait/0` switch.

¹⁵TODO: Details of *Tholds*, is this a return parameter?

First, if `AnteHoldsTODO == []`, we continue with `setup_lt_notground_fv/13` with the increased time interval. Otherwise we follow *AnteHoldsTODO* Literals.

NEXT: What does `find_min_range_ground(Atom, PN, Tholds, O2)` do? Probably: Inspect Literal starting from `Tholds`.

setup_lt_wait_var(FV, FVL, LitData, ToDoAnte, AnteHolds, TMin, THolds, ConseRId, PV, Delay, Id, IdTerm)

Probably: We know *AnteHolds* is ok between *TMin* and *THolds*. *LitData* is the current literal that has been analyzed. FV are the free variables in the Literal, FVL is the list of instantiations of FV that have been dealt with, for which this setup is not responsible at all. Called by `setup_lt_notgroundldefault/12` (same arguments) which is only called as last call in `setup_lt_notground_fv/13`.

Although we know *AnteHolds* is true starting from *TMin*, we also know that no *LitData* literal not having *FVL* instantiation is true starting from the current handled time.

We need to check that *TMin* is not relevant! Indeed it is not: see `update_activity_time1/2` for the `wait_var` term. *TMin* is only used to characterize the *Removed* argument for identifying the current activity. So, ... TODO: later remove the unnecessary *TMin* argument from `wait_var` terms.

TODO

- Really nail down the meaning of `wait_var`, also at what stage are what values for *TMin* and *THolds* set.
- Will `wait_vars` become invalidated? Inspect `get_new_tholds/15`.
- For a `wait_var` entry, can a result be found that extends *AnteHolds* starting at *TMin*? Probably: it could be that some rule has not fired yet, giving a new *LitData* instantiation. But please create an example.

5 wait entries

`wait` entries use an additional time *T* that determines whether the entry should be handled in the current step:

```
update_activity_times1(ResultTime) :-
    (    get_wait1(T, Activity, Ref),
      cmp_lt(T, ResultTime),
      retract_wait1(Ref),
      update_activity_time(Activity, T),
      fail
    ;    true
  ).
```

¹⁶When analysing the source:Be aware that in `get_new_tholds/15` the `Atom` in `[ds_lh(lit(Atom,PN),Id1,IdTerm1)|AnteHoldsTODO]` is ground, and has nothing to do with FV and FVL.

17

Times involved in `wait` entries:

TStart from what time do we need new true LitData

TEnd up to what time does *AnteHolds* hold, for `wait_var`

THandle At or before what HandledTime should `wait` entry be updated.

TFirstResult First time that consequent could become true, see `activity_min_result_time/2`.

As we cannot efficiently select those wait facts with $T < \text{HandledTime}$, we could inspect all wait entries and calculate relevant times each time.

5.1 `wait_var`

TStart: We should or could, start inspecting the wait entry as soon as new values are available from that time. Alternatively we could wait till $TStart + g$. So, **THandle** could be **TStart**, or $TStart + g$.

The first time the consequent should become true: $TStart + g + e$

At this time, we store `wait_var(Id, IdTerm, HT, TMin, FV, FVL, LitData, ToDoAnte, AnteHolds)`

As shown elsewhere, *TMin* is not relevant, *HT* is the time up to which we looked at *LitData*, *FVL*, and up to where there is no value available and from where the value is unknown/blank.

6 So far

Try documenting the whole data structure that describes the state of each leadsto rule first. All invariants, the understanding of having every possible outcome of a leadsto rule represented.

At what stage the HandledTime invariant is. Understanding the `get_new_tholds`, the reverse is on purpose as that is part of the invariant, having a partial instantiation left to right.

After that, try understanding the cleanup efforts of `wait_var`.

7 "Syntax" of Leadsto Specifications

Copied from `syntax.txt`:

The leadsto input syntax is prolog syntax, but with the following added/changed operator definitions.
(For input of leadsto specs in prolog, the `:redefinition` is awkward. I do a push/pop operator call for reading)

```
op(150, xfx, :),
op(700, xfy, <),
op(700, xfy, <=),
op(700, xfy, =<),
```

¹⁷ResultTime is the current/new, just updated HandledTime.

```

        op(700, xfy, >),
        op(700, xfy, >=)
    ]).

```

Currently, only the top level terms are described. I am working on syntax (+minimal explanation of semantics) of the top level terms, but especially the sub terms.

sub terms:

% VAR:PLPCE:

VAR, in principle, a prolog term, although Uppercase atoms are allowed. Quotes around atoms are allowed.

TODO: are unquoted uppercase functors allowed?

e.g. P(a) TODO: what are further restrictions and

interpretations of PLPCE terms TODO: junk this stupid

name "PLPCE".

% start_time(PLPCE)

% end_time(PLPCE)

% global_lambda(PLPCE)

TODO: why those qterms?

% qterm(cwa(X)) cwa_node

% qterm(external(X)) external_node

% qterm(X) ... other_node

% display(_,_)

% display_number_range(_,_,_,_)

% periodic(Vars, Range, Period:PLPCE, Formula) is_list(Vars)

 * handle_interval(Vars, Range, Formula, Root, Son, Extra)

% periodic(ST, ET, Period:PLPCE, Formula)

 * handle_interval([], range(ST, ET), Formula, Root, Son, Extra),

% periodic(Vars, ST, ET, Period:PLPCE, Formula)

 * handle_interval(Vars, range(ST, ET), Formula, Root, Son, Extra)

% interval(Vars, ST, ET, Formula)

 * handle_interval(Vars, Range, Formula, Root, _Son, Extra)

% interval(Vars, ST, ET, Formula)

 * handle_interval(Vars, range(ST, ET), Formula, Root, _Son, Extra)

% interval(ST, ET, Formula)

 * handle_interval([], range(ST, ET), Formula, Root, _Son, Extra)

% leadsto(AnteFormula, ConseFormula, Delay)

 * handle_leadsto1(Root, AnteFormula, ConseFormula, Delay, Extra, _Son)

% leadsto(Vars, AnteFormula, ConseFormula, Delay)

```

        * handle_leadstol(Root, AnteFormula, ConseFormula, Delay, Extra, Son)
% specification(_)
        * IGNORED
% content(C)
        * TODO? assertz(dyn_content(C))
% denotes(Header, Formula)
        * term_to_formula_node(Formula, FormulaNode, Extra),
        new(PN, property_def_node(@off)),
        send(PN, fill_header, Header),
        send(PN, son, FormulaNode),
        send(Root, son, PN)
% (sort_element(SortName:PLPCE, Term):- member(Term2, List) with Term==Term2
        * test_sort_def(SortName, List, Extra),
        ensure_sort_son(Root, SortName1, SNode),
        add_sort_contents(SNode, List).
% constant(Name, Value)
        * check_constant(Name, Value),
        send(Root, son, new(N, constant_def_node)),
        send(N, fill_header, Name),
        send(N, fill_value, Value)
% sortdef(SortName:PLPCE, Objs)
        * test_sort_def(SortName, Objs, Extra),
        send(Root, son, new(SN, sort_node)),
        send(SN, change_gui_prop, sort_name, SortName1),
        add_sort_contents(SN, Objs).
% cwa(PLPCE)
% model(PLPCE)
% [specification_element]
        * generic node

```

8 TODO

8.1 waitvar, waitfired

get_wait1/3 is often called with all variables leading to calls of current_key/1 and looping over all wait_var entries for each first_possible_activity_result/1 call.

8.1.1 Debugging, verifying waitvar FV,FVL aspects

I changed check_isolated_fire_rest. We need to test it for cases where there are variables. Make sure that check_isolated_fire_rest/8 is fired in some example, sometimes with excluded FV,FVL, sometimes not.

Follow FV, FVL with FV == []

Can a wait_var or wait_fired fire in case of FVL=[]?

do_default_cwa_isolated_no_trace: \+ find_atom_trace: we need to have the neg value in the range! Probably first generate an example where we have an atom trace but cwa holds in some

range and a rule should fire there.

Is it an option to produce a simpler less efficient algorithm that we can use to validate the result?

8.2 Solving bugje-nondeterminism

check_isolated_fire_rest/9 added FV, FVL so that we do not fire FV out of FVL.

To trust this more, I would like to understand the gc aspects. Let us follow wait_fired

8.3 analyse case

We have a leadsto rule

```
if a and b then (e=0, f=1, g=1, h=1) c
interval(a, 0, 1)
interval(b, 0, 1)
```

After we have fired the rule once, we have wait_var entries for a and b. At a later stage we should remove those entries as they have FV == [] and FVL == [[]] so they should never lead to new results. But the programming was sloppy. A number of aspects: the wait_var should respect the FVL restriction. It seems to do that but not when we replace the entry by an updated one in get_new_tholds/15. The problem could also be that we do not make the wait_var start from a new place.

Let us follow the second step where we have updated HandledTime to the start of the c derived interval. We encounter wait_var(ds_lh(b), ..).get O2 = fail(1, _)

In setup_lt_notground_fv/13 I see: assert_debug(cmp_ge(THolds, HT)). That should be the case for all setup_lt_wait_var, set_wait_var/14 calls, so check rm_gc_wait_vars1/8.

Still, I would prefer formulating the semantics of wait_var over this skimming and avoiding.

Where do we set wait_var? We have recorded dyn_wait(HT, Activity) and dyn_wait_var(Id1, IdTerm, HT, TMin, FV, FVL, LitData, ToDoAnte, AnteHolds, THolds, wait_var is set by set_wait1_var/14 and removed by retract_wait1/2 or retract_wait1/1 (erase/1). set_wait1_var/14 only called by set_wait_var/14. set_wait_var/14 called by setup_lt_wait_var/12 which adds HandledTime as an argument. set_wait_var/14 also called by rm_gc_wait_var1/8.

Let us check wait_fired too. It calls rm_gc_wait_wait_vars/8. This seems to be ok as the instantiated antecedent has no uninstantiated variables left.

TODO: Do we assure that all wait records are removed before a next run?

Some doubt about at what time to hang wait entries. Do we add e to the value because the result can only come after a delay or do we need to fire rules as soon as possible, i.e. when a new non-blank value can be available? In that case, we could also add g, as a rule could only results at e + g after the current HandledTime?

The question is: how do we define HandledTime?

1. All values up to HandledTime have been derived?
2. Or all values up to HandledTime have been derived and all results of applying rules on intervals up to HandledTime are in?

I will need to rethink this, I probably chose the first option. But the second one seems to be possible too. Let me put some time into checking the second case although it seems hard to determine the next `HandledTime`.

Possibly there could be more options even. The `wait` entries should not be linked to a new `HandledTime` value?

The `wait` entries should be defined independent from `HandledTime`, as they are now.

`wait_var` is linked to one `LitData` entry. Of `FV` and `FVL` we know that no value is known before `TStart`. We know that `AnteHolds` holds up to `THolds`. No other time should be needed. CLEAR! NO DOUBT.

We need to check this. Especially look at updating `wait` entries.

We store `wait` entries that depend on changes in the future. We can also calculate from what time a result could be expected. So after a "step" that checks whether rules may fire given `wait` entries and then updating `wait` entries, we determine `HandledTime`, the earliest time any of the `wait` entries could produce a result. Then the next step follows.

Issues:

- The actual firing? When analysing a `wait` entry, a rule could fire. Should we fire immediately? Should we wait? Does it matter? It seems that the sooner we decide to fire a rule, the more we know and the further we could extend "true" intervals. The firing could well change values before the new `HandledTime`, or could it? Probably not! It is the definition of the new `HandledTime`: no `wait` entry can produce results before the new `HandledTime`. Why do we gather `dyn_schedule_fire/3` facts and fire them in one batch using `handle_fired/0`? In the current implementation `handle_fired/0` is called at the beginning of `handled_time_step/1` and once at the end, after the last `handled_time_step/1`. If `update_activity_times/1` only schedules rules to fire, then before doing the next `update_activity_times/1`, all scheduled firings belonging to the previous `uat/1` should have fired. But still, why not fire immediately? It could be that we somehow depend on a known state of the future over one `wait` batch? Let us postpone this.